





Clustering

Dr. Holger Fröhlich
SS 2016

OMICs-Data (e.g. gene expression)



Preprocessing



Biological Samples

High-Throughput Measurements
(NGS, Microarrays)

n Samples

p Genes / Transcripts

	A	B	C	D	E	F	G	H	I	J	K
1	t1	t2	t3	t4	t5	t6	t7	t8	t9	t10	t11
2	YAL022C	3.8518	1.4433	4.9007	1.5214	0.41538	-1.6848	-0.04249	-2.3566	-3.0103	-0.06228
3	YAL023C	-2.4834	-1.8468	-1.8762	-4.0699	-1.8187	-4.0882	-2.8394	-2.4193	-2.2955	-0.83477
4	YAL036C	-0.8302	-0.03962	1.1755	0.051383	-0.15474	0.3497	-2.3273	-0.31484	1.1737	-0.04292
5	YAL037W	0.95071	-0.34414	1.7837	1.375	0.14011	0.90682	-0.31151	0.65269	0.025381	0.23071
6	YAL041V	-2.9395	-2.7089	-3.2279	-6.6413	-2.2626	-4.3677	-4.3092	-2.9473	-2.8837	-1.6481
7	YAL042W	0.88046	1.4121	0.70091	1.5236	0.536	1.6386	1.4624	0.93259	1.4878	0.62257
8	YAL043C	-2.3066	-1.9819	-2.6073	-4.7824	-2.2387	-4.2046	-3.0809	-2.895	-2.3703	-0.88976
9	YAL043C	0.59475	0.74273	1.3922	1.359	0.99807	1.1322	1.2846	1.2975	1.0213	0.49802
10	YAL044C	0.13819	0.51711	0.26241	0.6641	0.34171	1.5442	1.0322	1.0142	0.84372	0.34719
11	YAL045C	-0.89036	-2.7762	-2.94	-2.832	-3.1648	-4.5947	-3.3343	-4.1272	-4.5737	-2.8244
12	YAL054C	-0.61552	-0.8198	-0.29818	-0.84141	-0.75644	-1.1779	-1.1553	-0.6179	-0.60902	0.4404
13	YAL063C	-0.61299	0.055744	-0.16914	-0.73895	-0.1452	-0.39563	0.644	0.10609	0.21114	-0.57642
14	YAR007C	-1.1401	-0.68046	-0.17562	-0.93679	-0.26384	0.10037	-0.69386	-0.20379	-0.8507	-0.4815
15	YAR008W	-0.89949	-0.32658	-0.45516	0.28005	-0.68723	-0.03708	-0.17731	0.031561	-0.41564	-0.55937
16	YAR009C	0.37513	0.57632	-0.4956	0.27061	-0.28603	0.40515	-0.53192	-0.65724	0.45596	0.034053
17	YAR050V	-0.03397	0.62255	-2.586	0.40751	-0.69945	2.1786	-0.29562	-2.1935	3.1602	0.14045
18	YBL007C	0.40274	0.40606	0.15697	0.63259	1.1127	0.8843	1.0171	0.85515	0.99862	0.37357
19	YBL008W	0.062519	-0.33747	-1.0013	-0.95188	-0.81554	-0.54217	0.25262	0.35317	0.16779	-0.35719
20	YBL017C	-0.41402	0.16599	-0.08462	-0.08169	0.045784	0.82145	0.54198	-0.24443	1.0108	-0.24005
21	YBL029W	0.75188	-1.2895	1.2904	2.3651	0.89355	0.63978	-0.29606	0.97384	-0.78895	0.37852
22	YBL030C	-0.10457	-0.89876	0.55978	0.25046	0.37137	0.34062	0.2064	-0.0273	-0.73219	0.28842
23	YBL038W	-0.41717	-0.14104	-0.01792	-0.4331	-0.05168	-0.34599	-0.09384	-0.19882	-0.25844	0.2349
24	YBL039C	-1.5146	-1.7394	-1.4437	-3.001	-1.1918	-2.1556	-2.2566	-2.0463	-1.7803	-0.61371
25	YBL079W	1.711	2.2484	2.4589	4.1205	1.5702	3.4163	3.6512	2.2029	2.3956	0.44653

...

p = 20.000 – 50.000

n << p

patients: n <= few 100

Cells / cell lines: n <= 5 (per biol. condition)


Cluster Analysis

	A	B	C	D	E	F	G	H	I	J	K
1		t1	t2	t3	t4	t5	t6	t7	t8	t9	t10
2	YAL022C	3.8518	1.4433	4.9007	1.5214	0.41538	-1.6848	-0.04249	-2.3566	-3.0103	-0.06228
3	YAL023C	-2.4834	-1.8468	-1.8762	-4.0699	-1.8187	-4.0882	-2.8394	-2.4193	-2.2955	-0.83477
4	YAL026C	-0.83282	-0.03952	1.1755	0.061383	-0.15474	0.3497	2.3273	-0.31494	1.1737	-0.04292
5	YAL037W	0.95071	-0.34414	1.7837	1.375	0.14011	0.90682	-0.31151	0.65269	0.025381	0.23071
6	YAL041W	-2.9395	-2.7089	-3.2279	-5.6413	-2.2626	-4.3877	-4.3092	-2.9473	-2.8837	-1.6481
7	YAL042W	0.86046	1.4121	0.70091	1.5236	0.536	1.8386	1.4524	0.93259	1.4878	0.62257
8	YAL043C	-2.3066	-1.9819	-2.6073	-4.7824	-2.2387	-4.2046	-3.0809	-2.895	-2.3703	-0.88976
9	YAL043C	0.59475	0.74273	1.3922	1.359	0.99807	1.1322	1.2846	1.2975	1.0213	0.49802
10	YAL044C	0.13819	0.51711	0.28241	0.6641	0.34171	1.5442	1.0322	1.0142	0.84372	0.34719
11	YAL045C	-0.85836	-2.7762	-2.94	-2.832	-3.1648	-4.5947	-3.3343	-4.1272	-4.5737	-2.8244
12	YAL054C	-0.61552	-0.8198	-0.29818	-0.84141	-0.75644	-1.1779	-1.1553	-0.6179	-0.60902	0.4404
13	YAL063C	-0.61299	0.055744	-0.16914	-0.73895	-0.1452	-0.39563	0.644	0.10609	0.21114	-0.57642
14	YAR007C	-1.1401	-0.68046	-0.17562	-0.93679	-0.26384	0.10037	-0.69386	-0.20379	-0.8507	-0.4815
15	YAR008W	-0.89949	-0.32658	-0.45516	0.28005	-0.68723	-0.03708	-0.17731	0.031561	-0.41564	-0.55937
16	YAR009C	0.37513	0.57632	-0.4956	0.27061	-0.28803	0.40515	-0.53192	-0.65724	0.45586	0.034053
17	YAR050W	-0.03397	0.62255	-2.586	0.40751	-0.69945	2.1786	-0.29562	-2.1935	3.1602	0.14045
18	YBL007C	0.40774	0.40606	0.15697	0.63259	1.1127	0.8843	1.0171	0.85515	0.99982	0.37357
19	YBL008W	0.060519	-0.33747	-1.0013	-0.95188	-0.81554	-0.54217	0.25262	0.39317	0.16779	-0.35719
20	YBL017C	-0.41402	0.16599	-0.08462	-0.08169	0.045784	0.82145	0.54198	-0.24443	1.0108	-0.24005
21	YBL029W	0.75188	-1.2895	1.2904	2.3651	0.89355	0.63978	-0.29606	0.97384	-0.78985	0.37852
22	YBL030C	-0.10457	-0.89976	0.55978	0.25046	0.37137	0.34062	0.2064	-0.0273	-0.73219	0.28942
23	YBL038W	-0.41717	-0.14104	-0.01782	-0.4331	-0.06168	-0.34599	-0.09384	-0.18862	-0.25844	0.2349
24	YBL039C	-1.5146	-1.7394	-1.4437	-3.001	-1.1918	-2.1556	-2.2566	-2.0463	-1.7803	-0.61371
25	YBL079W	1.711	2.2494	2.4589	4.1205	1.5702	3.4163	3.5612	2.2029	2.3956	0.44653

- Find groups (clusters) of genes with similar expressions profile (*co-expressed* genes)
- No true grouping known: **unsupervised** learning problem

Geometric Interpretation


- Genes = points
- Find clusters of similar points
- No true clustering known
- How to find a „good“ clustering?




Homogeneity and separation principle

- **Homogeneity:** Elements within a cluster are close to each other
- **Separation:** Elements in different clusters are further apart from each other

good



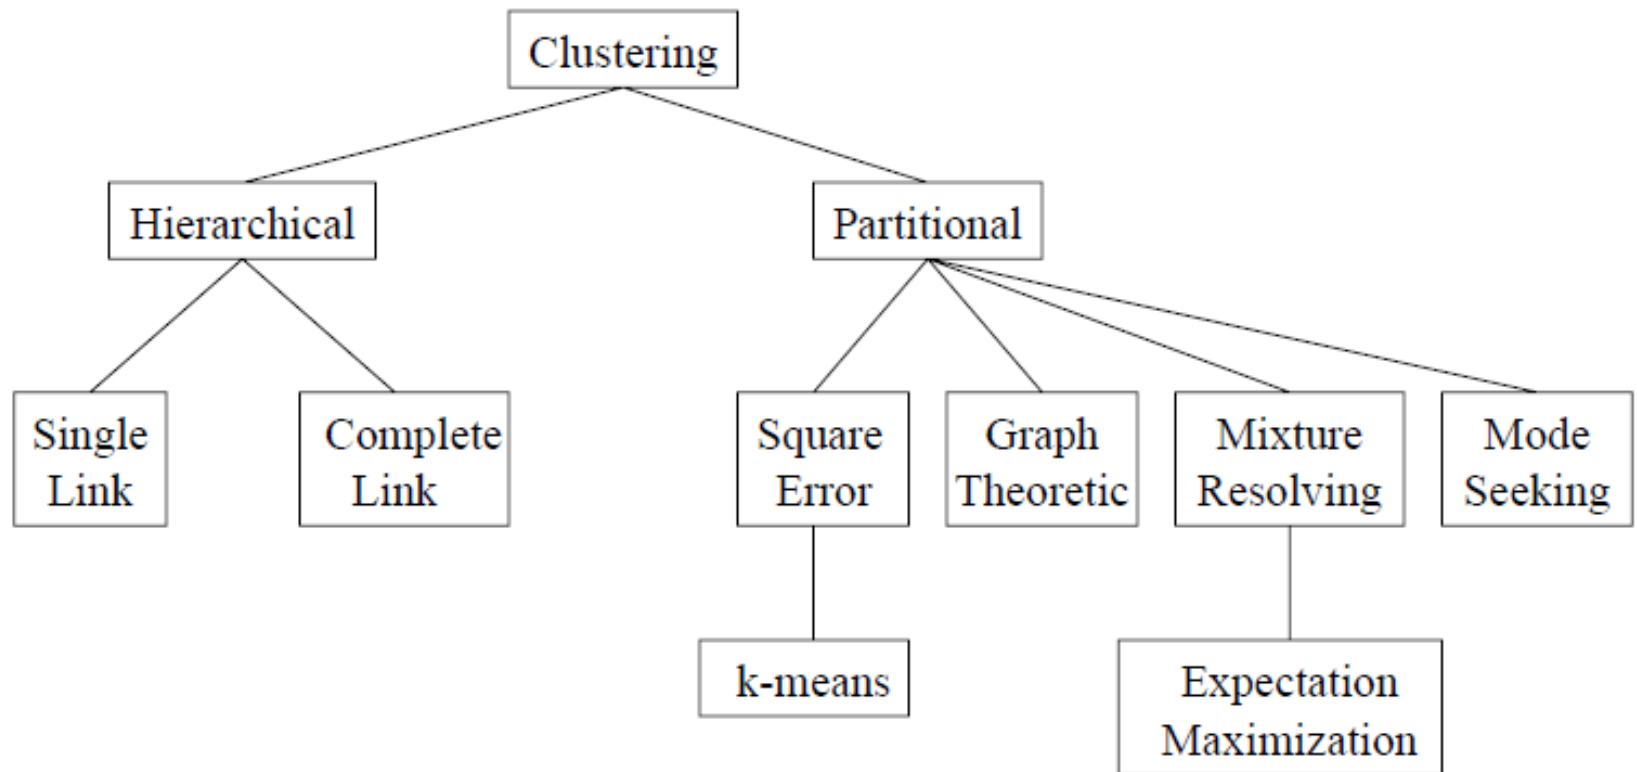
bad



Clustering algorithms


- No mathematically unique way of defining these two general criteria
 - *ill defined problem*
- Many clustering algorithms following different approaches
 - Success of individual algorithm is data dependent
- Difficult to validate clustering results
 - Clustering is (usually) *unsupervised*: there is no objective ground truth or „true“ clustering

Taxonomy of clustering techniques (Jain et al., 1999)



Hierarchical vs. Partitioning Clustering

- **Partitioning:** Given goal function and fixed number k of clusters
 - Divide objects in k partitions, such that goal function is optimized (e.g. mean distance to cluster centers)
- **Hierarchical:** agglomerative oder divisive
 - **Now:** agglomerative hierarch. clustering



Proximity Measures

- Agglomerative hier. Clustering depends on:
 - How we measure the proximities of objects (genes)
 - How we measure the similarity of clusters (and thus fuse most similar ones)
- For metric data (such as gene expression) distances can be used for objects.
- Minkowski / p-norm distance / metric:

$$d(x, y) = \left(\sum_{i=1}^n |x[i] - y[i]|^p \right)^{1/p} = \|x - y\|_p$$


- Special cases:
 - Euclidian distance (p=2)
 - Manhattan distance (p=1)

Distance measures

- Pearson correlation „distance“ (not a metric):

$$d(x, y) = 1 - \rho_{x,y} = 1 - \frac{\sum_i (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_i (x_i - \bar{x})^2 \sum_i (y_i - \bar{y})^2}}$$


- Often used for time series clustering



- Euclidean distance: green more similar to red
- Pearson distance: green more similar to black


Cluster fusion: How similar are two clusters?

Single Linkage




$$D_{SL}(A, B) = \min_{\substack{a \in A \\ b \in B}} d(a, b)$$

Complete Linkage




$$D_{CL}(A, B) = \max_{\substack{a \in A \\ b \in B}} d(a, b)$$

Average Linkage



$$D_{AL}(A, B) = \frac{1}{|A| |B|} \sum_{\substack{a \in A \\ b \in B}} d(a, b)$$

Ward Criterion



$$D_{Ward}(A, B) = \frac{d(\bar{a}, \bar{b})}{1/|A| + 1/|B|}$$



Algorithm for agglomerative hierarchical Clustering

- Inputs:
 - n = Number of objects to cluster (genes)
 - D = $n \times n$ distance matrix
- Output: hierarchical clustering (*dendrogramm*)
- Example: **Complete Linkage Clustering**

	A	B	C	D	E
A	0	1	2	3	5
B	1	0	3	4	4
C	2	3	0	1	4
D	3	4	1	0	2
E	5	4	4	2	0




	U	C	D	E
U	0	3	4	5
C	3	0	1	4
D	4	1	0	2
E	5	4	2	0




Next Steps





	U	C	D	E
U	0	3	4	5
C	3	0	1	4
D	4	1	0	2
E	5	4	2	0



	U	V	E
U	0	4	5
V	4	0	4
E	5	3	0



	W	E
W	0	5
E	5	0



Complexity

- $n - 1$ cluster merging steps
- In each of these steps $O(n^2)$ possibilities, to join clusters
- **(naive) overall complexity: $O(n^3)$**
 - Using priority queues: $O(n^2 \log n)$
- For Single and Complete Linkage: improvement to $O(n^2)$ possible
- Agglomerative hierarchical clustering is relatively computational expensive
 - Typical applications: up to few thousand objects

Lance & Willilams Formula

- Observation: algorithm only updates distances of new cluster $U=\{A,B\}$ to any existing cluster C
- General formula for update (**Lance & Williams, 1966**):

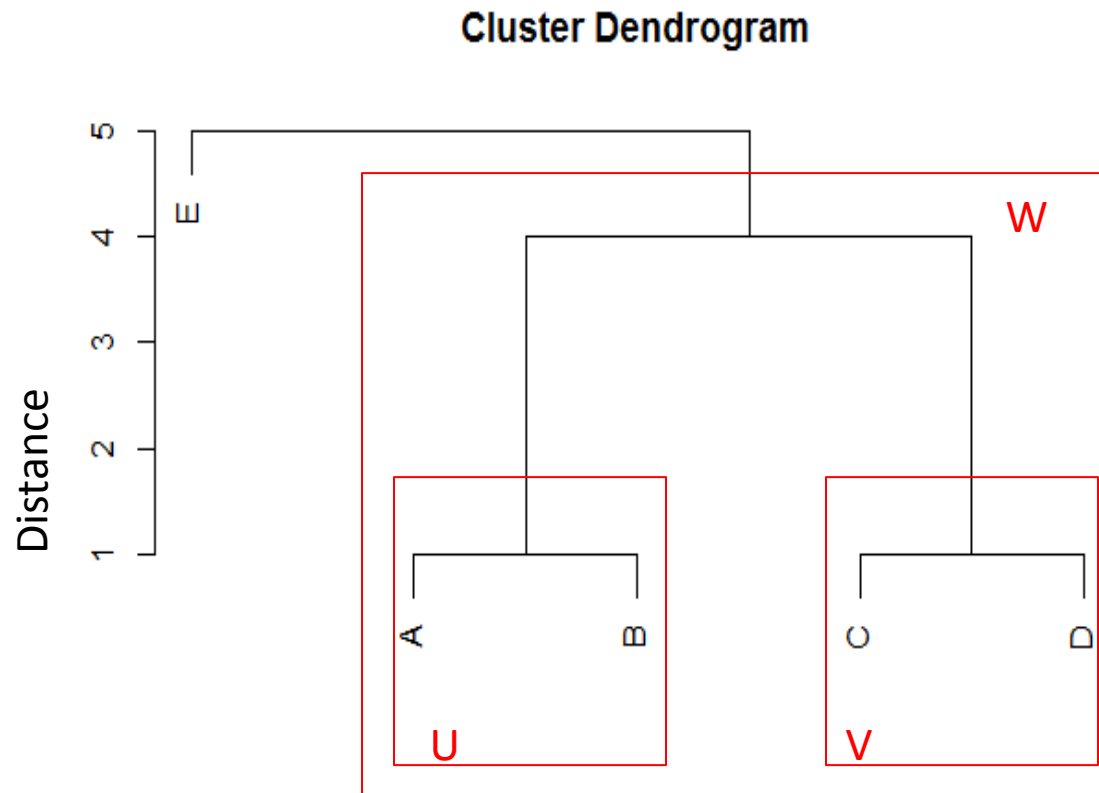
$$D(U, C) = \alpha_1 D(A, C) + \alpha_2 D(B, C) + \beta D(A, B) + \gamma |D(A, C) - D(B, C)|$$

UPGMA

Methode	α_1	α_2	β	γ
Single L.	$\frac{1}{2}$	$\frac{1}{2}$	0	$-\frac{1}{2}$
Complete L.	$\frac{1}{2}$	$\frac{1}{2}$	0	$\frac{1}{2}$
Weigthed group average L.	$\frac{1}{2}$	$\frac{1}{2}$	0	0
Unweighted group Av. L.	$\frac{ A }{ A + B }$	$\frac{ B }{ A + B }$	0	0
Ward	$\frac{ A }{ A + B + C }$	$\frac{ B }{ A + B + C }$	$-\frac{ C }{ A + B + C }$	0


Algorithm Output: The Dendrogram

- Memorize during algorithm execution:
 - Which clusters were merged
 - Which distances they have




Reading clustering dendrograms

- Clustering dendrogram is NOT unique!
- **Example:** subtree in red box could also be drawn on the right side!




Influence of cluster similarity measure on clustering result

Single linkage




Complete linkage




- Single linkage: long, chain-like clusters
- Complete linkage: small, ellipsoidal clusters
- Ward: spherical, isotropic clusters

Example: Colon Cancer Data (Alon, 1999)


Single L.




Complete L.



Average L.

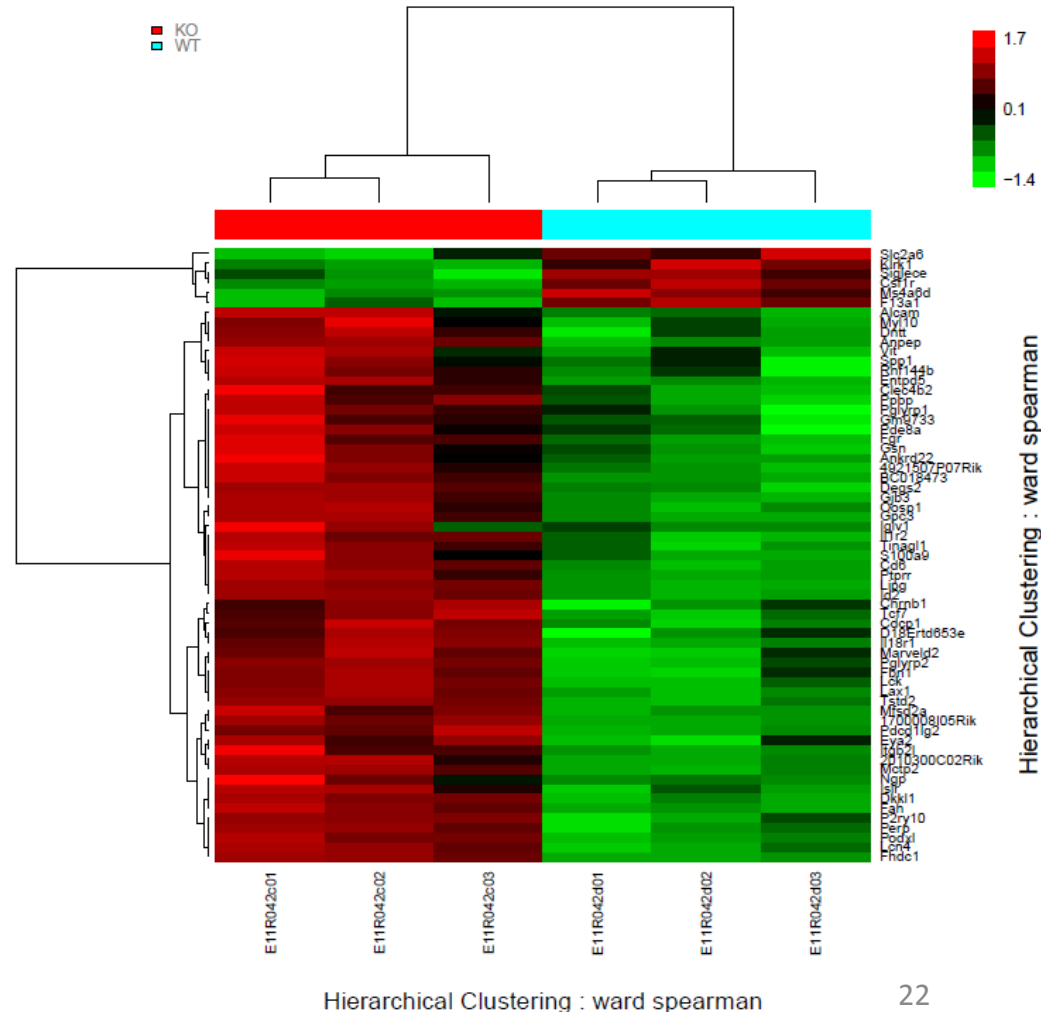


Ward



Popular Application: Heatmaps

- Represent whole data matrix with false colors
 - Color code = expression level
- Group rows (genes) and columns (samples) via hierarchical clustering
 - Can see groups of related samples + genes at once



Features of Hierarchical Clustering

- Advantages:
 - Visual data analysis: no pre-specified number of clusters
 - User can specify a cut point in the hierarchy, e.g. where intercluster distance exceeds some threshold
 - Organizes the clusters in a hierarchical way: dendrograms
 - BUT: dendrogram is not unique
 - Clustered heatmaps as application
- Drawbacks:
 - larger datasets: long run time, messy dendrograms

Partitioning Algorithms: K-means

- One of the most prominent **partitioning** algorithms
- **Idea:** represent data in terms of K clusters, each summarized via prototype a μ_k
- Each data point is assigned to one cluster
- Minimize distortion:

$$\sum_{k=1}^K \sum_{x_j \in C_k} \|x_j - \mu_k\|^2$$


K-Means Clustering: Lloyd Algorithm

1. **K-means(k)**
2. Arbitrarily assign the k cluster centers
3. **while** the cluster centers keep changing
4. Assign each data point to the cluster C_i corresponding to the **closest** cluster representative (center) ($1 \leq i \leq k$)
5. After the assignment of all data points, compute new cluster representatives as


Euclidean distance

$$\mu_k = \frac{1}{|C_k|} \sum_{x_j \in C_k} x_j$$


Example: initialization




Example: cluster assignment




Example: new cluster centers




Example: cluster re-assignment




Example: new cluster centers




Example: cluster re-assignment



Example: final cluster centers





Example: final cluster assignment



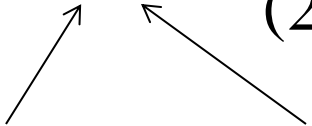
Notes on k-means

- Algorithm usually converges very fast
 - complexity: $O(n * k * d * i)$ with n = #samples; d = #variables; i = #iterations
- **BUT:** may lead to suboptimal solutions
- Clustering depends on initial conditions
 - Solution: repeat l times
- K-means can only detect spherical clusters!
- Hard assignment to clusters
 - Small shifts of a data point can flip cluster membership

Example: Mixture of 3 Gaussians



Multivariate Normal Distributions

$$N(\mathbf{x} | \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{(2\pi |\boldsymbol{\Sigma}|)^{d/2}} e^{-\frac{1}{2}(\mathbf{x}-\boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1}(\mathbf{x}-\boldsymbol{\mu})}$$


Mean vector
(d dimensions)



Covariance matrix (d x d)

Maximum likelihood estimates:

$$\hat{\boldsymbol{\mu}} = \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i$$

$$\hat{\boldsymbol{\Sigma}} = \frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i - \hat{\boldsymbol{\mu}})(\mathbf{x}_i - \hat{\boldsymbol{\mu}})^T$$

Multivariate Normal Distributions - Example




Gaussian Mixtures Models (GMMs)

- We consider our data to be drawn from a *mixture* of multivariate normal distributions:

$$p(x) = \sum_{k=1}^K \pi_k N(x | \mu_k, \Sigma_k) \text{ with } \sum_k \pi_k = 1, \pi_k \in [0,1]$$

- Interpretation of data generating process:
 - First pick a cluster (component) with probability π_k
 - Then draw a sample \mathbf{x}_i from that component
- Each data point is generated by one of K multivariate normal distributions


Gaussian Mixture Models: Schematic View



- We formally describe cluster membership of each data point via an indicator variable

GMMs (cont'd)

- Problem: true cluster membership of each data point unknown
- ➔ indicator variables are **hidden / latent**
- We would like to make inference on $z_i = (z_{i1}, \dots, z_{iK})$ given observed data



Likelihood

- We suppose data points to be drawn iid
- **Complete likelihood** of observed and unobserved variables:

$$p(\{x_i\}, \{z_i\} | \mu_1, \dots, \mu_K, \Sigma_1, \dots, \Sigma_K) = \prod_{i=1}^n \prod_{k=1}^K p(x_i | \mu_k, \Sigma_k, z_{ik} = 1) \Pr(z_{ik} = 1)$$

- **Problem:** direct maximization w.r.t. parameters AND unobserved variables not possible

Expectation Maximization (EM) algorithm – initialization

- Start with some cluster assignment
- Estimate parameters of each Gaussian via ML

$$\hat{\mu}_k = \frac{\sum_{i=1}^n z_{ik} x_i}{\sum_{i=1}^n z_{ik}}$$
$$\hat{\Sigma}_k = \frac{\sum_{i=1}^n z_{ik} (x_i - \hat{\mu}_k)(x_i - \hat{\mu}_k)^T}{\sum_{i=1}^n z_{ik}}$$
$$\hat{\pi}_k = \frac{\sum_{i=1}^n z_{ik}}{n}$$

Expectation Maximization (EM) algorithm – E-step

- Given parameters of each Gaussian: compute *expected* cluster assignment of each data point:

$$E[z_{ik}] = \Pr(z_{ik} = 1 \mid x_i, \mu_1, \dots, \mu_K, \Sigma_1, \dots, \Sigma_K) * 1 \\ + 0 * \Pr(z_{ik} = 0 \mid x_i, \mu_1, \dots, \mu_K, \Sigma_1, \dots, \Sigma_K)$$

$$\stackrel{\text{Bayeslaw}}{=} \frac{p(x_i \mid \mu_1, \dots, \mu_K, \Sigma_1, \dots, \Sigma_K, z_{ik} = 1) \Pr(z_{ik} = 1)}{\sum_{k=1}^K p(x_i \mid \mu_1, \dots, \mu_K, \Sigma_1, \dots, \Sigma_K, z_{ik} = 1) \Pr(z_{ik} = 1)}$$

Expectation Maximization (EM) algorithm – M-step

- Given: expected cluster assignments of data points
- Recompute ML estimates for parameters for each Gaussian

$$\hat{\mu}_k = \frac{\sum_{i=1}^n E[z_{ik}] x_i}{\sum_{i=1}^n E[z_{ik}]}$$


$$\hat{\Sigma}_k = \frac{\sum_{i=1}^n E[z_{ik}] (x_i - \hat{\mu}_k)(x_i - \hat{\mu}_k)^T}{\sum_{i=1}^n E[z_{ik}]}$$

$$\hat{\pi}_k = \frac{\sum_{i=1}^n E[z_{ik}]}{n}$$


Complete EM Algorithm

- Initialize cluster assignment (e.g. random or via k-means)
- Iterate until convergence (complete likelihood does not increase significantly):
 - E-step
 - M-step
- Algorithm is guaranteed to increase the complete likelihood in each iteration
- BUT: may get stuck into local optima
 - Sensitive to initialization


Example: Diabetes dataset (8 variables) – PCA plot




Initialization (random cluster assignment)




E-Step




M-Step




E-Step (2)



M-Step (2)



After convergence



EM Algorithm in general

- Let Θ denote the vector of all parameters and let Θ^t denote its current estimate.
- E-step: find expected value of complete log-likelihood

$$Q(\Theta | \Theta^t) := E_{p(z|x, \Theta^t)} [\log p(x, z | \Theta)] := \sum_{k=1}^K \log p(x, z = k | \Theta) P(z = k | x, \Theta^t)$$

- M-step: find the parameters that maximize Q:

$$\Theta^{t+1} = \arg \max_{\Theta} Q(\Theta | \Theta^t)$$

Q-function for GMMs

$$\begin{aligned} Q(\Theta | \Theta^t) &= \sum_{i=1}^n \sum_{k=1}^K \log p(x_i, z_{ik} = 1 | \Theta) \Pr(z_{ik} = 1 | x, \Theta^t) \\ &= \sum_{i=1}^n \sum_{k=1}^K \log p(x_i, z_{ik} = 1 | \Theta) \Pr(z_{ik} = 1 | x, \Theta^t) \\ &= \sum_{i=1}^n \sum_{k=1}^K (\log N(x_i | \mu_k, \Sigma_k, z_{ik} = 1) + \log \pi_k) \Pr(z_{ik} = 1 | x_i, \mu_1, \dots, \mu_K, \Sigma_1, \dots, \Sigma_K) \end{aligned}$$

Normal distribution density for point i

Prior probability for point i to belong to cluster k

Expected cluster membership for point i (slide 43)

Features of GMMs

- Advantages:
 - Probabilistic cluster assignment, clusters may overlap
 - Algorithm can detect ellipsoidal clusters of different sizes
 - GMMs are **model based**: model may be used to assign future data points to the most likely cluster or to impute missing values
- Disadvantages:
 - need to know number of clusters / mixture components K in advance
 - BUT: existant statistical heuristics (\rightarrow model selection) to estimate K from data
 - Slow for large amounts of data

Selecting the number of clusters k

- For GMMs we can use the so-called model selection criteria, e.g. Bayesian Information Criterion (BIC) to determine a good number of clusters from data:

$$BIC = -\log\text{--likelihood} + 0.5 * \log(n) * n_{par}$$

- **Rational (informal):** the more clusters we have, the more parameter we need to estimate effectively from data, i.e. the GMM model gets more and more complex

Selecting the number of clusters k (cont'd)

- **Problem with overly complex model:**
 - overfitting – other data drawn from the same distribution may not be explained well
 - GMM forms clusters, which are of minor information and do not help to interpret the data
- **Okkams razor principle:** Try to find a model, which is as simple as possible to explain your data sufficiently
- ➔ Clustering is a way to reduce data complexity. If too many clusters, nothing is won

Bayesian Information Criterion (BIC)

$$BIC = -\log\text{--likelihood} + 0.5 * \log(n) * npar$$

data fit

Complexity penalty

- $npar$ = number of parameters in the model to be estimated
 - Depends on number of clusters
- BIC balances fit to the data and model complexity
- Heuristic approach!

Using BIC in practice

1. Define a set of cluster numbers K
2. for each k in K :
 - a. Run GMM clustering
 - b. Determine BIC
3. Select clustering with the lowest BIC

Clustering Validity

- How can we check the quality of a clustering?
 - GMMs: log-likelihood (depends on k), BIC
 - k-means: distortion (depends on k)
- Problem: values are not on a normalized scale, can only be used in a relative sense
- Visual inspection
 - Plot the data itself (difficult for > 2 dimensions)
 - Plot distance structure (\rightarrow clustering silhouettes)
- Clustering indices to measure validity

Cluster Silhouettes (Rousseeuw, 1987)

- For each observation i assigned to cluster C the silhouette $s(i)$ is defined as:

$$s(i) := \begin{cases} 0 & D(i, C) = 0 \\ \frac{D(i, B) - D(i, C)}{\max(D(i, C), D(i, B))} & \text{otherwise} \end{cases} \in [-1, 1]$$

$$D(i, C) := \frac{1}{|C|} \sum_{a \in C} d(a, i)$$


$$D(i, B) := \min_{C' \neq C} D(i, C')$$

- Interpretation:
 - $s(i)$ close to 1: i lies in cluster C
 - $s(i) = 0$: i lies between two clusters
 - $s(i) < 0$: i is close to B than to C

Examples: good and bad clustering

Silhouette plot of pam(x = ruspini, k = 4)


n = 75



Average silhouette width : 0.74


Silhouette plot of (x = pr4\$clustering, dist = dist(ruspini))

n = 75



Average silhouette width : -0.08

Selecting the number of clusters




Clustering High Dimensional Samples

- **So far:** main focus on clustering of features
- How about clustering of samples?
- **Example:** Can we identify patient sub-populations from molecular data?
- Problem: far more features than samples

Clustering in High Dimensions

- Data points are sparsely distributed in a high dimensional space
 - Biology: only few features (< 5%) expected to show significant differences between samples
 - (Euclidean) distance dominated by noise features
- $$d(x, y) = \left(\sum_{i=1}^n |x[i] - y[i]|^2 \right)^{1/2}$$
 - All data points (patients) become almost equally distant from each other
 - We are clustering noise!
- Consequence: Leaving out or adding a few samples could drastically change clustering (**statistical instability**)



Clustering in High Dimensions

1. Reduce features
 - Pre-filter features
 - Use clustering algorithm with in-built selection of most relevant features
2. Look for a pattern that remains statistically stable, even if sample set changes slightly
 - Consensus over different clusterings

Example for Pre-filtering: Verhaak et al., Cancer Cell, 2010

Data: gene expression profiles of 206 patients with Glioblastoma Multiforme (brain tumor), 3 different technical platforms

Prefiltering of most variable genes:

1. High correlation across platforms → 9,255 genes
2. High variability (MAD) on each platform → 1,903 genes
3. Exclusion of genes with large differences in MAD across platforms → 1,740 genes

Consensus average linkage clustering (Monti et al., 2003) shows 4 cancer sub-types

- Motivation: address statistical instability

Consensus Clustering (Monti et al., Machine Learning, 2003)

Idea: assess clustering stability via sub-sampling

1. Sample $p\%$ (default: 80%) of the data points without replacement
2. Optional: do the same for the features
3. Run base clustering algorithm (e.g. k-means, average linkage)
4. Repeat H times

Main question: how to form consensus out of H clusterings?

Define connectivity matrix for clustering h :

$$M^{(h)}(i, j) = \begin{cases} 1 & \text{if items } i \text{ and } j \text{ belong to the same cluster,} \\ 0 & \text{otherwise.} \end{cases}$$

Let $I^{(h)}$ be a $N \times N$ matrix indicating the presence of i and j :

$$I^{(h)} = \begin{cases} 1 & i \text{ and } j \text{ in subsample } h \\ 0 & \text{otherwise} \end{cases}$$

Consensus Clustering

Consensus matrix is defined as properly normalized sum of all connectivity matrices:

$$\mathcal{M}(i, j) = \frac{\sum_h M^{(h)}(i, j)}{\sum_h I^{(h)}(i, j)}$$

Normalization takes into account whether both, i and j , are present

Observations:

- M is symmetric
- M has values in $[0,1]$ where 1 means perfect consensus
- M may be viewed as a similarity measure between items
- Reordering rows and columns in M according to true clustering yields a block-diagonal matrix

Consequence: Final clustering can be achieved via hierarchical clustering using $\mathbf{1} - \mathbf{M}$ as distance matrix

Summary Statistics

Consensus matrix provides information about

- Stability of overall clustering
- Stability of individual clusters
- Cluster representatives and outliers

Cluster k consensus:

$$m(k) = \frac{1}{N_k(N_k - 1)/2} \sum_{\substack{i, j \in I_k \\ i < j}} \mathcal{M}(i, j)$$

- Measures average frequency, how often objects in cluster k group together

Consensus item for cluster k:

$$m_i(k) = \frac{1}{N_k - 1\{e_i \in I_k\}} \sum_{\substack{j \in I_k \\ j \neq i}} \mathcal{M}(i, j)$$

- Measures average consensus of item e_i with all other items in cluster k
- **Representative**: item with highest consensus
- **Outliers**: opposite

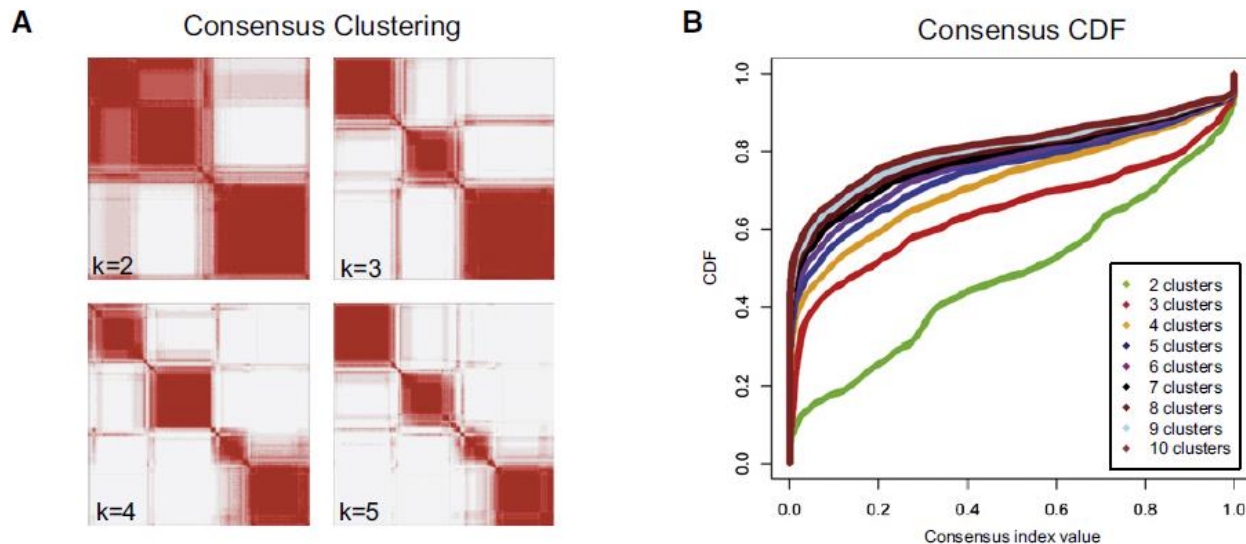
Determining the Number of Clusters

Summary statistics depend on chosen number K of clusters: How to find a good K ?

Idea: consider the distribution of values in consensus matrix

- Stronger skew to 1 \rightarrow higher stability

Plot empirical CDF of this distribution for different K



Determining the Number of Clusters

Consider the area under the CDF curves


$$A(K) = \sum_{i=2}^m [x_i - x_{i-1}] \text{CDF}(x_i)$$

Observation (example): area increases significantly from 3 to 4 clusters and then stabilizes

- More clusters cannot be detected stably

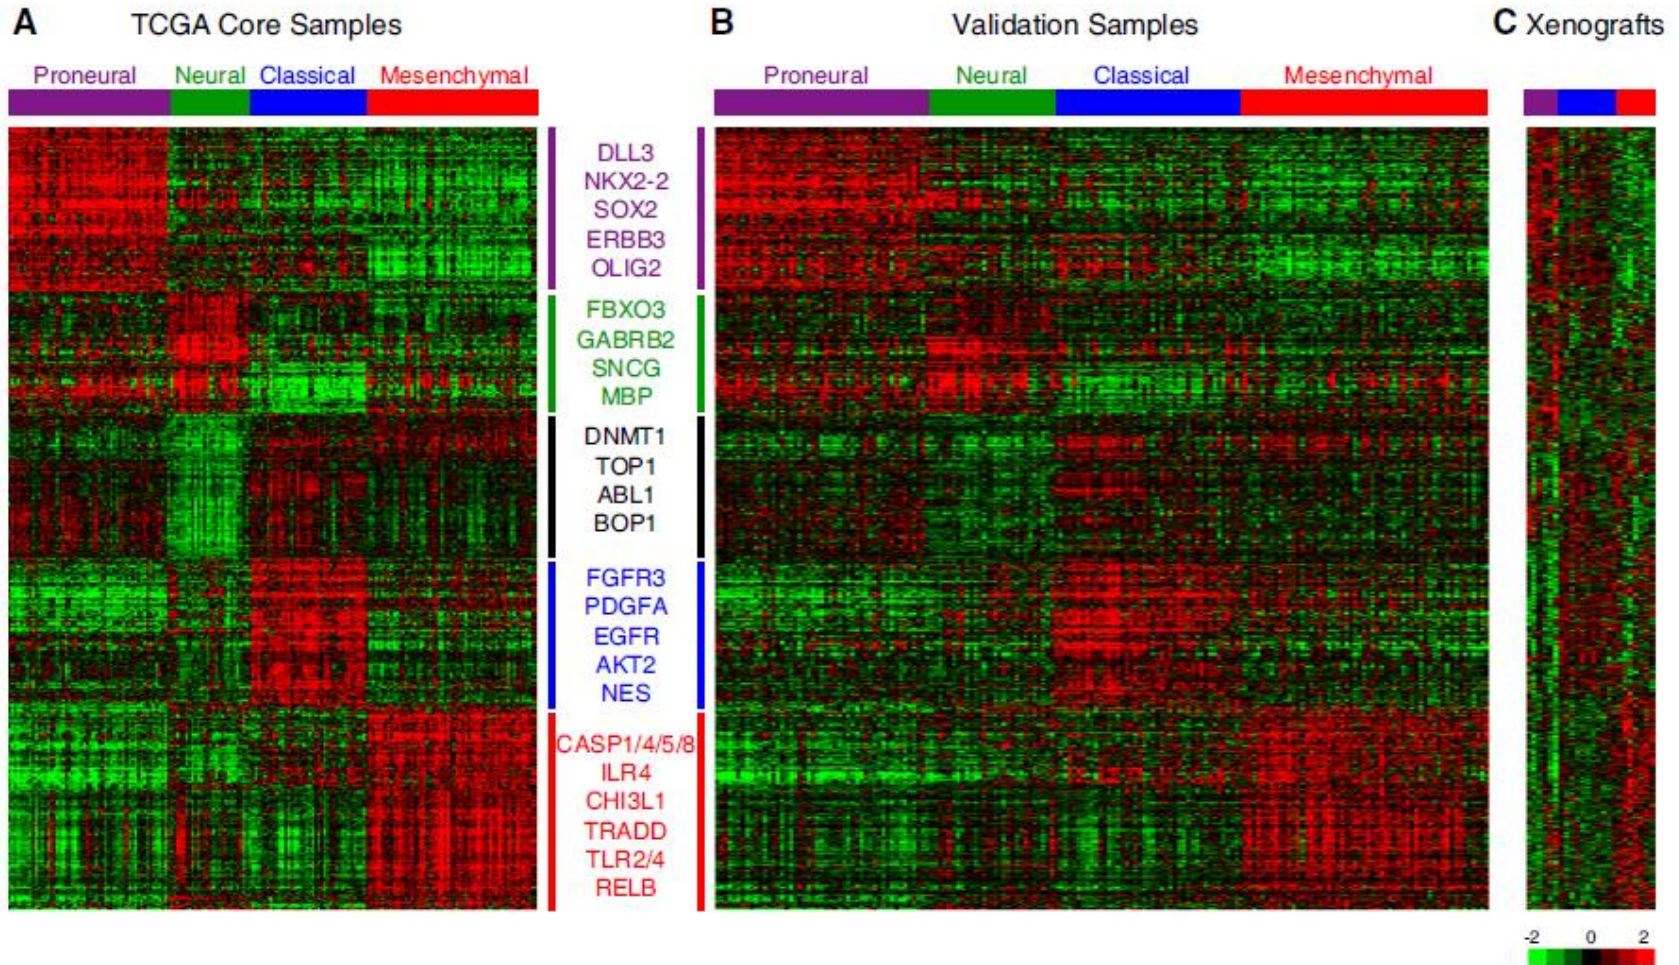
Thus: plot K against proportional change of $A(K)$

$$\Delta(K) = \begin{cases} A(K) & \text{if } K = 2 \\ \frac{A(K+1) - A(K)}{A(K)} & \text{if } K > 2, \end{cases}$$



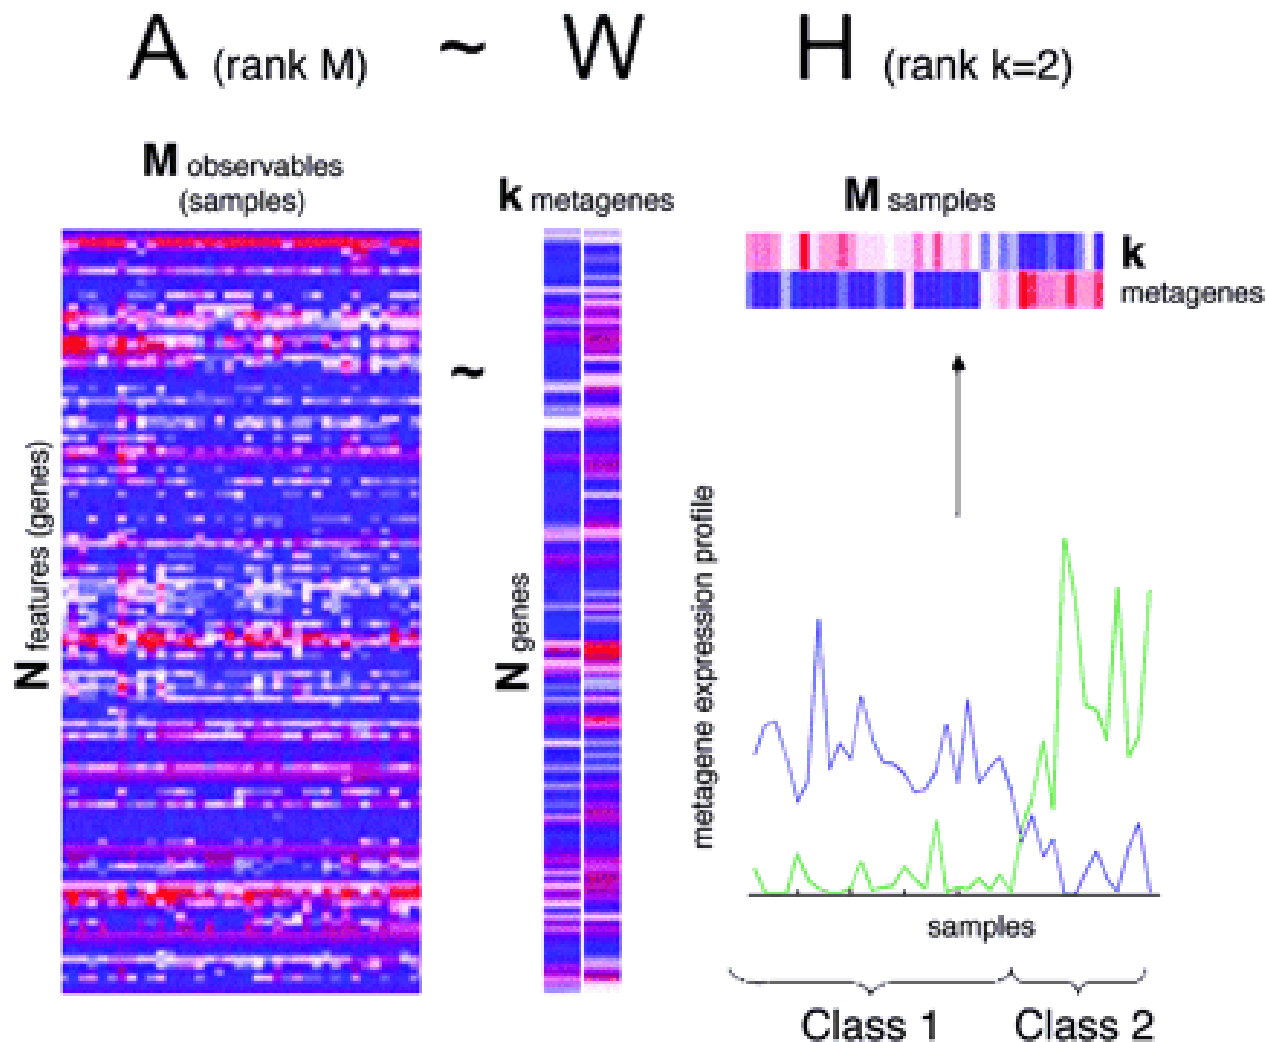
(c)

Consensus Clustering Result



Alternative Strategy: Non-Negative Matrix Factorization (Brunet et al., PNAS, 2004)

Can different leukemia sub-types be identified from gene expression profiles?




Non-negative Matrix Factorization (Lee & Seung, Nature, 1999)

Consider $p \times N$ data matrix X with non-negative entries

Goal: find low rank approximation of X such that

$$X \approx WH$$

- W : $p \times r$, **non-negative**
- H : $r \times N$, **non-negative**
- $r \ll \max(N, p)$



Approach: $\arg \min_{W, H} \|X - WH\|_F$ subject to $W, H \geq 0$

- Non-convex
- Not unique

An Algorithm for Solving NMF (Lee & Seung, 1999)

Objective: minimize divergence between X and WH

$$D(X \| WH) = \sum_{i,j} \left(x_{ij} \log \frac{x_{ij}}{w_{ij}h_{ij}} - x_{ij} + w_{ij}h_{ij} \right)$$

Find local minimum by alternating two steps:

$$w_{ik} \leftarrow w_{ik} \frac{\sum_{j=1}^P h_{kj} x_{ij} / (WH)_{ij}}{\sum_{j=1}^P h_{kj}}$$

$$h_{kj} \leftarrow h_{kj} \frac{\sum_{i=1}^N w_{ik} x_{ij} / (WH)_{ij}}{\sum_{i=1}^N w_{ik}}$$

Several other algorithmic variants known


NMF Consensus Clustering

NMF procedure is prone to random initialization

Idea in Brunet et al. (2004): perform a number of runs and analyze consensus matrix, like in Consensus Clustering

- Hierarchical clustering based on consensus matrix
- Choose appropriate number of clusters based on silhouette index or something similar

Example (from Brunet et al.): clustering of leukemia patients based on gene expression data



Summary

- Clustering is an **exploratory** technique
 - Examples:
 - identification of co-expressed genes
 - Identification of patient sub-populations
 - No prediction
 - Hard to validate
 - No claims about statistical significance of
 - Overall clustering structure
 - Differences between defined clusters
 - Statistical stability
- Hierarchical clustering, k-means, GMMs, NMF as examples
- Other frequently used methods in Bioinformatics
 - K-medoids (PAM)
 - Self Organizing Maps (SOMs)
- Consensus clustering as a means to address statistical instability


More Applications in Bioinformatics

- Grouping of homologous sequences into gene families (evolutionary biology)
- Phylogenetic tree reconstruction
- Multiple sequence alignment
- Biological or medical image analysis

What you should know and being able to apply

- What is the purpose of clustering?
- Which clustering techniques exist (those which were covered here) and how do they work in principle?
- What are the pros and cons of the individual clustering algorithms discussed here? Which kind of clusters can they detect?
- How can we determine the quality of a clustering and the number of clusters?

Literature



- Monti, S., Tamayo, P., Mesirov, J. & Golub, T (2003). Consensus clustering: A resampling-based method for class discovery and visualization of gene expression microarray data. Machine Learning 52, 91–118.
- Daniel D. Lee and H. Sebastian Seung (1999). Learning the parts of objects by non-negative matrix factorization. Nature 401 (6755): 788–791
- Brunet, Tamayo, Golub, Mesirov (2004). Metagenes and molecular pattern discovery using matrix factorization. PNAS 101(12):4164–4169
- Verhaak, R. et al. (2010). Integrated Genomic Analysis Identifies Clinically Relevant Subtypes of Glioblastoma Characterized by Abnormalities in PDGFRA, IDH1, EGFR, and NF1. Cancer Cell 17, 98–110