

Chapter 4: Dimensionality Reduction

Jun.-Prof. Dr.-Ing. Thomas Schultz

URL: <http://cg.cs.uni-bonn.de>

E-Mail: schultz@cs.uni-bonn.de

Office: Friedrich-Ebert-Allee 144, 53113 Bonn

November 22, 2016

Dimensionality Reduction

Goal: Embed high-dimensional data into a low-dimensional space, in a way that important structure is preserved

We will cover four methods:


- **PCA:** Brief recap, you should already know this
- **Kernel PCA:** Deal with nonlinear structures
- **MDS:** Finds a low-dimensional embedding based on distances alone
- **ISOMAP:** Allows us to “unroll” nonlinear manifolds

Section 4.1:

(Kernel) Principal Component Analysis

PCA – the general idea


- PCA finds an orthogonal basis that best represents a given data set.



The sum of squared distances from the x' axis is minimized.

PCA – the general idea

- PCA finds an orthogonal basis that best represents a given data set.




3D point set in
standard basis

PCA finds a best approximating plane
(again, in terms of $\sum \text{distances}^2$)

Managing high-dimensional data


- Example: Data base of face scans (3D geometry + texture)




- 10,000 points in each scan
- x, y, z, R, G, B – 6 numbers for each point
- Thus, each scan is a $10,000 \times 6 = \mathbf{60,000}$ -dimensional vector!

Managing high-dimensional data

- How to find interesting axes in this 60,000-dimensional space?
 - axes that measure age, gender, etc...
 - There is hope: the faces are likely to be governed by a small set of parameters (much fewer than 60,000...)



age axis



gender axis

Goals of Principal Component Analysis


- Given input $\mathbf{x}_i \in \mathbb{R}^p$, $i = 1, 2, \dots, n$, learn mapping from \mathbf{x} to $\mathbf{y} = \mathbf{U}_k^T(\mathbf{x} - \bar{\mathbf{x}})$
 - **Approximation:** $\mathbf{U}_k \in \mathbb{R}^{p \times k}$ with $k \leq p$
 - Select \mathbf{U}_k such that the approximation error $\sum_i \left\| (\mathbf{I} - \mathbf{U}_k \mathbf{U}_k^T)(\mathbf{x}_i - \bar{\mathbf{x}}) \right\|^2$ is minimal
 - Select \mathbf{U}_k so that the variance in the projected data remains maximal
 - **Decorrelation:**
 - Coefficients of \mathbf{y} are uncorrelated
 - If $k = p$, pure rotation, no approximation

Algorithm: PCA

- **Input:** $\mathbf{x}_i \in \mathbb{R}^p, i = 1, 2, \dots, n$
- **Output:** $U_k \in \mathbb{R}^{p \times k}$ with $k \leq p$ such that $\mathbf{y} = U_k^T (\mathbf{x} - \bar{\mathbf{x}})$ with $\bar{\mathbf{x}} = \frac{1}{n} \sum_i \mathbf{x}_i$
- **Algorithm:**
 - **Center** the points, $\mathbf{z}_i = \mathbf{x}_i - \bar{\mathbf{x}}$
 - Form $p \times n$ **centered data matrix** $Z = [\mathbf{z}_1, \mathbf{z}_2, \dots, \mathbf{z}_n]$
 - Form $p \times p$ **scatter matrix** $S = ZZ^T$
 - Compute spectral decomposition $S = U\Lambda U^T$
 - Sort coefficients of Λ in decreasing order
 - Form U_k from k leading columns of U
- I assume you have seen the derivation previously,

Principal Components

- Eigenvectors that correspond to **big** eigenvalues are the directions in which the data has strong components (= large variance).
- If the eigenvalues are more or less the same, there is no main direction.




There is no preferred direction.

S looks like this:

$$\begin{pmatrix} \lambda & \\ & \lambda \end{pmatrix}$$

Any vector is an eigenvector



There is a clear preferred direction.


S looks like this:

$$U \begin{pmatrix} \lambda & \\ & \mu \end{pmatrix} U^T$$

μ is close to zero, much smaller than λ .


Example: Biological Shape Analysis

- **Data:** CT images of 48 mouse mandibles
- **Scientific question:** Impact of factors such as evolutionary history, diet, or geography on skeleton shape



Linear Shape Spaces



- **Align** all shapes, compute a **mean shape**
- Express each shape as transformation of the mean
 - **Shape vector**: Concatenation of displacement vectors
- Perform PCA on shape vectors to identify **principal modes** of shape variation



Linear Shape Space: Example



- Normal component: Color
- Tangential component: Arrow glyph

PC1
(52%)




9.6

PC2
(21%)




8.7



Exploring Shape Spaces



Hermann et al.:
 „A Visual Analytics
 Approach to Study
 Anatomic
 Covariation.“
 PacificVis 2014



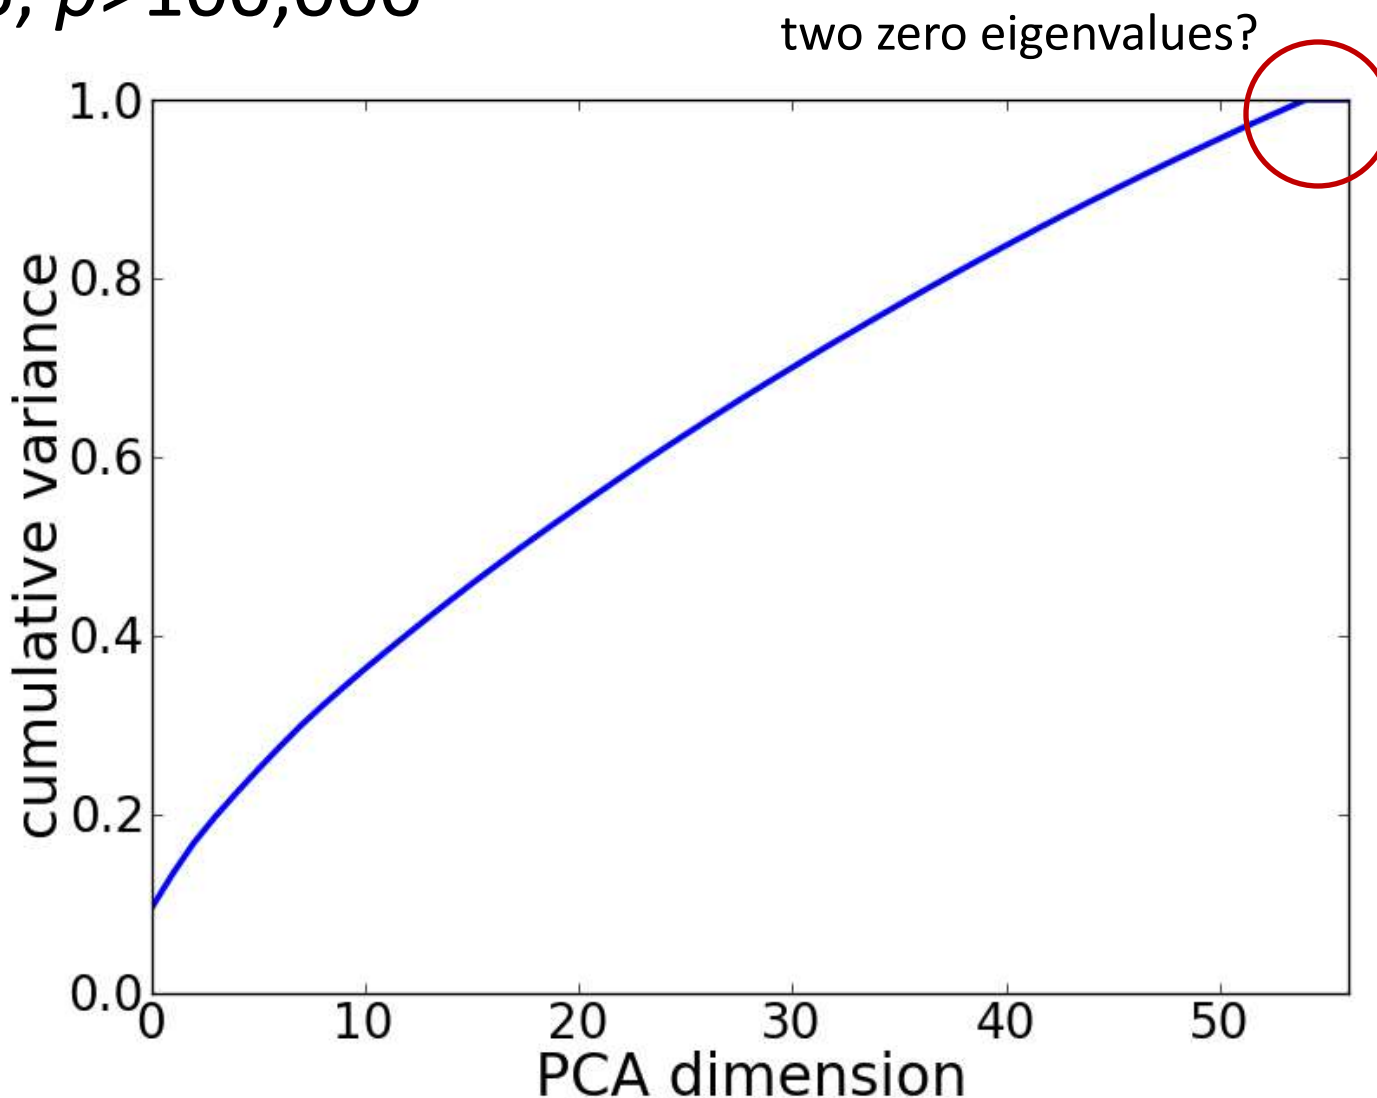
Group Analysis



Hermann et al.: „Accurate Interactive Visualization of Large Deformations and Variability in Biomedical Image Ensembles.“
IEEE SciVis 2015

Quiz: Interpreting PCA Result

$n=56$, $p>100,000$



What if $n < p$?

- **Maximum rank** of Z and S is $\min(n - 1, p)$
 - Linear dependence introduced by centering
 - It does not make sense to use $k > n - 1$
- All eigenvectors of S are in the span of \mathbf{z}_i :

$$\begin{aligned}\lambda_a \mathbf{u}_a &= S \mathbf{u}_a = \sum_i \mathbf{z}_i \mathbf{z}_i^T \mathbf{u}_a = \sum_i (\mathbf{z}_i^T \mathbf{u}_a) \mathbf{z}_i \\ \Rightarrow u_a &= \sum_i \frac{\mathbf{z}_i^T \mathbf{u}_a}{\lambda_a} \mathbf{z}_i =: \sum_i \alpha_i^a \mathbf{z}_i\end{aligned}$$

Efficient computation for $n < p$

- We can replace the $p \times p$ with an $n \times n$ eigenvalue problem
 - Based on $S\mathbf{u} = \lambda\mathbf{u} \Leftrightarrow \mathbf{z}_i^T S\mathbf{u} = \lambda \mathbf{z}_i^T \mathbf{u}$ for all i
$$\mathbf{z}_i^T S\mathbf{u}_a = \lambda_a \mathbf{z}_i^T \mathbf{u}_a$$
$$\mathbf{z}_i^T \sum_l \mathbf{z}_l \mathbf{z}_l^T \sum_j \alpha_j^a \mathbf{z}_j = \lambda_a \mathbf{z}_i^T \sum_j \alpha_j^a \mathbf{z}_j$$
$$\sum_{j,l} \alpha_j^a [\mathbf{z}_i^T \mathbf{z}_l][\mathbf{z}_l^T \mathbf{z}_j] = \lambda_a \sum_j \alpha_j^a [\mathbf{z}_i^T \mathbf{z}_j]$$
 - With $n \times n$ centered inner product („Gram“)
matrix $K_{ij}^c := [\mathbf{z}_i^T \mathbf{z}_j]$:
$$(K^c)^2 \boldsymbol{\alpha}^a = \lambda_a K^c \boldsymbol{\alpha}^a \Rightarrow K^c \boldsymbol{\alpha}^a = \lambda_a \boldsymbol{\alpha}^a$$

Proper Normalization and Projecting New Data

- **Normalization:** $\mathbf{u}_a^T \mathbf{u}_a = 1$ translates into

$$\sum_{i,j} \alpha_i^a \alpha_j^a [\mathbf{z}_i^T \mathbf{z}_j] = (\boldsymbol{\alpha}^a)^T \mathbf{K}^c \boldsymbol{\alpha}^a \\ = \lambda_a (\boldsymbol{\alpha}^a)^T \boldsymbol{\alpha}^a = 1$$

– Therefore:


$$\|\boldsymbol{\alpha}^a\| = \frac{1}{\sqrt{\lambda_a}}$$

- **Projection** of a new point \mathbf{x} :


$$\mathbf{u}_a^T (\mathbf{x} - \bar{\mathbf{x}}) = \sum_i \alpha_i^a \mathbf{z}_i^T (\mathbf{x} - \bar{\mathbf{x}})$$

Preserving Nonlinear Structures

- PCA works well for **linear structures**
 - Straight lines, planes, etc.
- Can we preserve **nonlinear / curved structure**?
- **Idea:** Nonlinearly map data to a higher-dimensional feature space, apply PCA there



$$(x_1, x_2) \mapsto (x_1, x_2, x_1^2 + x_2^2)$$



Kernel Trick

- We reformulated PCA so that it only makes use of the data \mathbf{x}_i within scalar products
- To apply PCA in feature space, we do not explicitly need the feature map $\mathbf{x} \rightarrow \phi(\mathbf{x})$, only a nonlinear kernel function $K(\mathbf{x}_i, \mathbf{x}_j) = \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j)$
- K should produce positive definite matrices
- Widely used examples (here without proofs):
 - Linear: $K(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i^T \mathbf{x}_j$
 - Polynomial of power p : $K(\mathbf{x}_i, \mathbf{x}_j) = (1 + \mathbf{x}_i^T \mathbf{x}_j)^p$
 - Gaussian (RBF): $K(\mathbf{x}_i, \mathbf{x}_j) = e^{-\gamma \|\mathbf{x}_i - \mathbf{x}_j\|^2}$

Centering in Feature Space

- We require the centered kernel matrix $K_{ij}^c = (\Phi_i - \bar{\Phi})^T (\Phi_j - \bar{\Phi})$, but evaluating the kernel function gives us $K_{ij} = K(\mathbf{x}_i, \mathbf{x}_j) = \Phi_i^T \Phi_j$
 - Applying K to centered data \mathbf{z}_i is *not* equivalent to centering feature vectors Φ_i

- Centering can be performed as follows:

$$\begin{aligned} K_{ij}^c &= (\Phi_i - \bar{\Phi})^T (\Phi_j - \bar{\Phi}) \\ &= \Phi_i^T \Phi_j - \bar{\Phi}^T \Phi_j - \Phi_i^T \bar{\Phi} + \bar{\Phi}^T \bar{\Phi} \end{aligned}$$

- $\bar{\Phi}^T \Phi_j = \frac{1}{n} \sum_i \Phi_i^T \Phi_j = \frac{1}{n} \sum_i K_{ij}$
- $\Phi_i^T \bar{\Phi} = \frac{1}{n} \sum_j \Phi_i^T \Phi_j = \frac{1}{n} \sum_j K_{ij}$
- $\bar{\Phi}^T \bar{\Phi} = \frac{1}{n^2} \sum_{i,j} \Phi_i^T \Phi_j = \frac{1}{n^2} \sum_{i,j} K_{ij}$

Notation: Centering in Feature Space

- Let $\mathbf{1} \in \mathbb{R}^n$ denote the column vector in which each coefficient equals 1.
 - Then, $H = I - \frac{1}{n} \mathbf{1}\mathbf{1}^T$ is a matrix with $\frac{n-1}{n}$ on the diagonal and $-\frac{1}{n}$ everywhere else
- Easy to verify that

$$K^c = HKH$$

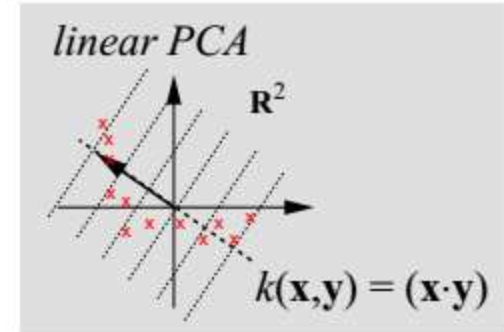
Algorithm: Kernel PCA

- **Input:** $\mathbf{x}_i \in \mathbb{R}^p, i = 1, 2, \dots, n$ and kernel function $K: \mathbb{R}^p \times \mathbb{R}^p \rightarrow \mathbb{R}$
- **Output:** $U_k \in \mathbb{R}^{p \times k}$ with $k \leq p$ and $D_k \in \mathbb{R}^{k \times k}$ such that $\mathbf{y} = D_k^{-1} U_k^T H \left(\mathbf{k} - \frac{1}{n} K \mathbf{1} \right)$ with $k_i = K(\mathbf{x}_i, \mathbf{x})$
- **Algorithm:**
 - Compute **kernel matrix** $K_{ij} = K(\mathbf{x}_i, \mathbf{x}_j)$
 - **Center kernel matrix** $K^c = HKH$
 - Compute **spectral decomposition** $K^c = U \Lambda U^T$
 - Sort coefficients of Λ in decreasing order
 - Form U_k from k leading columns of U
 - Form D_k from top-left $k \times k$ block of $\sqrt{\Lambda}$

Kernel PCA: Advantages and Disadvantages


- **Advantages** of Kernel PCA:

- Better reflects nonlinear structures
- Given suitable kernels, can be applied to more abstract objects or to unfold manifolds (see next sections)



- **Disadvantages** of Kernel PCA:

- Less interpretable: Principal modes are no longer a fixed combination of input variables
- Unlike with linear PCA, it is not easy to find a vector that corresponds to given Kernel PCA coordinates („pre-image problem“)

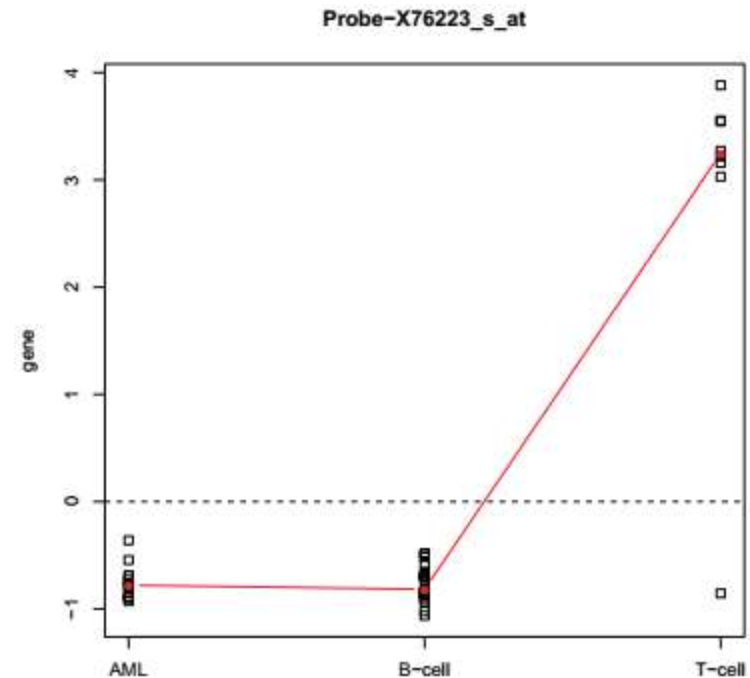
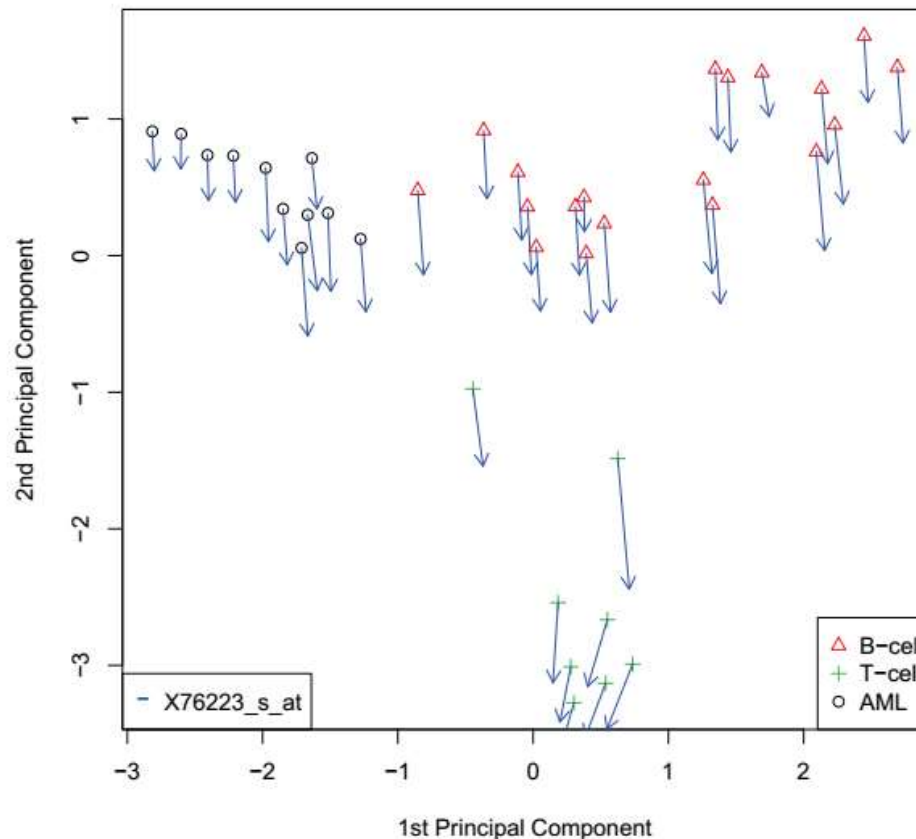


- **Advantage / Disadvantage:**

- Complexity grows with number of samples
- If needed, cluster as pre-process (e.g., k-means)

Example: Kernel PCA of Gene Expression Data

- Leukemia dataset with $n=19$, $p=3051$
- Kernel PCA with RBF kernel



Section 4.2:

Multidimensional Scaling

Multidimensional scaling (MDS)

Multidimensional scaling (MDS)

- A dimensionality reduction technique
- Maps pairwise distances to coordinates
- Provides model of non-geometric data
- Useful for
 - visualizing high-dimensional data
 - preprocessing data before clustering

Example: Multidimensional scaling (MDS)

Example: map of the US

Given a list of cities...

Chicago

Raleigh

Boston

Seattle

San Francisco

Austin

Orlando

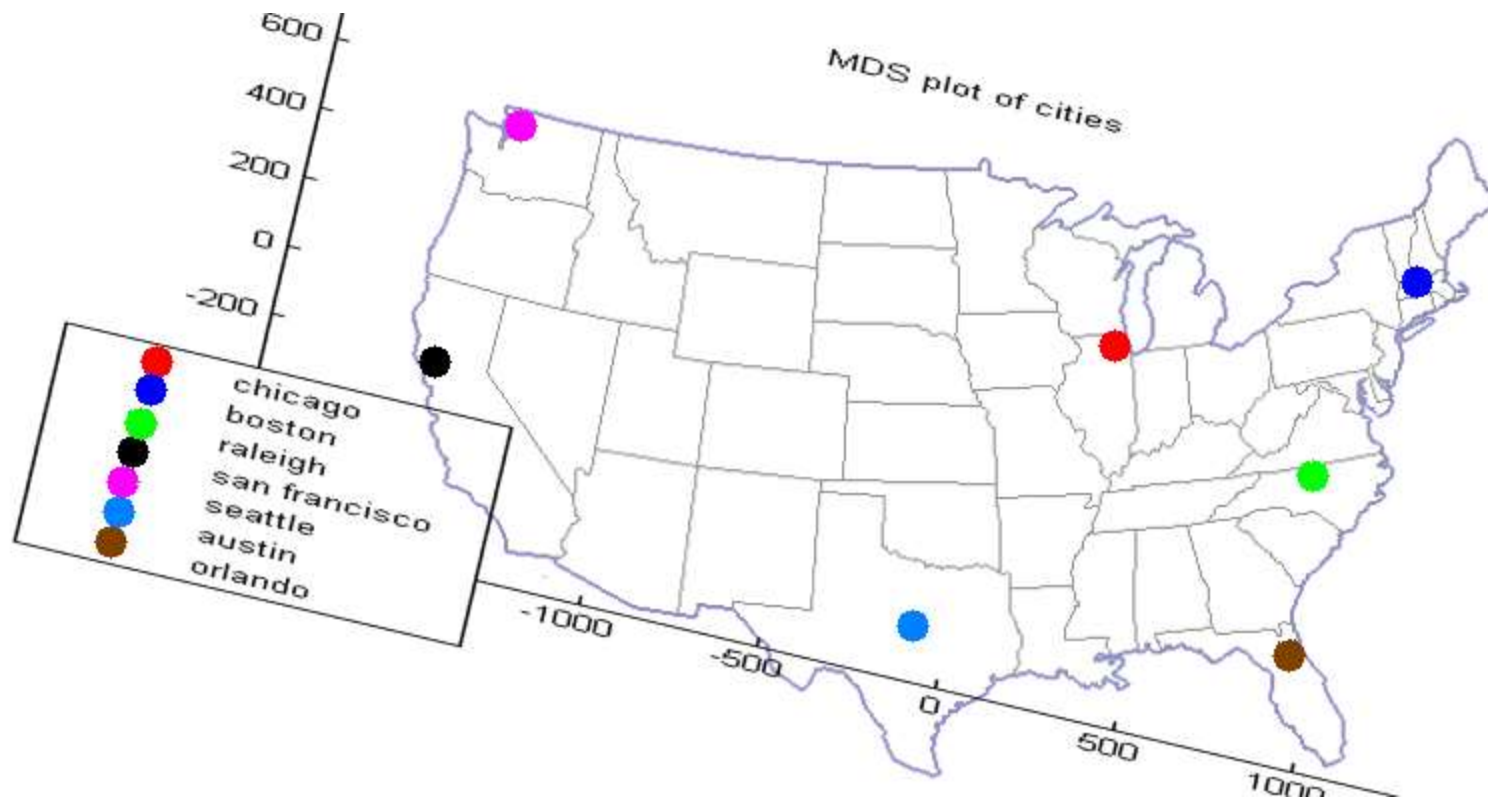
Example: Multidimensional scaling (MDS)

Knowing only the distances between them...

	Chicago	Raleigh	Boston	Seattle	San Francisco	Austin	Orlando
Chicago	0						
Raleigh	641	0					
Boston	851	608	0				
Seattle	1733	2363	2488	0			
San Francisco	1855	2406	2696	684	0		
Austin	972	1167	1691	1764	1495	0	
Orlando	994	520	1105	2565	2458	1015	0

Example: Multidimensional scaling (MDS)

Result: Given only the distances between points, MDS finds suitable coordinates for each point, of the specified dimension




Example with non-Euclidean Dissimilarity Matrix

- Given: *street* distances between 16 German cities (according to a route planner)

	B	HB	DD	D	EF	HH	H	KI	MD	MZ	M	P	SB	SN	S	WI
Berlin (B)	0	392	193	559	303	289	286	353	156	580	585	36	722	211	633	570
Bremen (HB)	392	0	472	287	334	124	131	208	253	476	749	363	575	224	629	466
Dresden (DD)	193	472	0	615	218	500	367	542	232	496	465	211	638	400	510	486
Düsseldorf (D)	559	287	615	0	400	411	277	495	419	216	611	529	298	511	404	201
Erfurt (EF)	303	334	218	400	0	362	219	458	235	289	414	274	431	475	339	279
Hamburg (HH)	289	124	500	411	362	0	157	97	280	528	775	283	670	109	655	518
Hannover (H)	286	131	367	277	219	157	0	248	147	384	632	256	526	252	512	374
Kiel (KI)	353	208	542	495	458	97	248	0	376	624	872	348	766	175	752	614
Magdeburg (MD)	156	253	232	419	235	280	147	376	0	465	519	125	607	309	567	455
Mainz (MZ)	580	476	496	216	289	528	384	624	465	0	430	544	146	621	210	14
München (M)	585	749	465	611	414	775	632	872	519	430	0	555	487	757	232	424
Potsdam (P)	36	363	211	529	274	283	256	348	125	544	555	0	693	207	604	541
Saarbrücken (SB)	722	575	638	298	431	670	526	766	607	146	487	693	0	771	262	159
Schwerin (SN)	211	224	400	511	475	109	252	175	309	621	757	207	771	0	756	619
Stuttgart (S)	633	629	510	404	339	655	512	752	567	210	232	604	262	756	0	220
Wiesbaden (WI)	570	466	486	201	279	518	374	614	455	14	424	541	159	619	220	0

MDS: Non-Euclidean Example

- MDS result



MDS: Non-Euclidean Example

- MDS result rotated and overlaid on a geographic map



Multidimensional scaling: Basic Concepts

- Input
 - Symmetric dissimilarity matrix, containing pair-wise dissimilarities $\delta_{ij} = \delta(x_i, x_j)$ between (p -dimensional) data samples x_1, \dots, x_n
 - Desired dimensionality k ($k < p$, often $k = 2, 3$)
- Output
 - Image of data in a Euclidean frame y_1, \dots, y_n such that pair-wise distances $d_{ij} = d(y_i, y_j)$ are best approximation to original dissimilarities $\delta(x_i, x_j)$ for k dimensions
- Embedding errors
 - Take into account $n(n - 1)/2$ errors between the individual distances
 - Have to be invariant against rigid transformations (i.e. translation, rotation, mirroring)

Find $y_i \in \mathbb{R}^k$ such that $d(y_i, y_j) \approx \delta_{ij}$

MDS: What is a good approximation?

- Possible error functionals (“Stress”):

$$J_{ee} = \frac{\sum_{i < j} (d_{ij} - \delta_{ij})^2}{\sum_{i < j} \delta_{ij}^2} \quad (\text{punishes large absolute deviations})$$

$$J_{ff} = \sum_{i < j} \left(\frac{d_{ij} - \delta_{ij}}{\delta_{ij}} \right)^2 \quad (\text{punishes large relative deviations})$$

$$J_{ef} = \frac{1}{\sum_{i < j} \delta_{ij}} \sum_{i < j} \frac{(d_{ij} - \delta_{ij})^2}{\delta_{ij}} \quad \text{Compromise between } J_{ee} \text{ and } J_{ff}$$

(Commonly L2 distance measure $d_{ij} = \|y_i - y_j\|_2$ is used here.)

MDS: What is a good approximation?


- Example global vs. local

$$J_{ee} = \frac{\sum_{i < j} (d_{ij} - \delta_{ij})^2}{\sum_{i < j} \delta_{ij}^2}$$

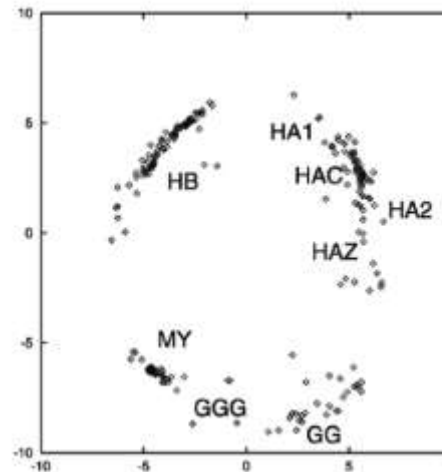
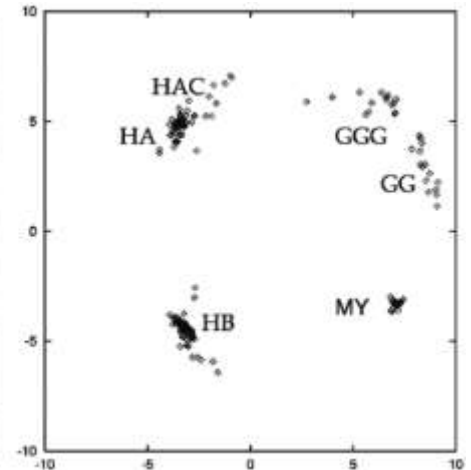
$$J_{ff} = \sum_{i < j} \left(\frac{d_{ij} - \delta_{ij}}{\delta_{ij}} \right)^2$$

$$J_{ef} = \frac{1}{\sum_{i < j} \delta_{ij}} \sum_{i < j} \frac{(d_{ij} - \delta_{ij})^2}{\delta_{ij}}$$

Similarity matrix

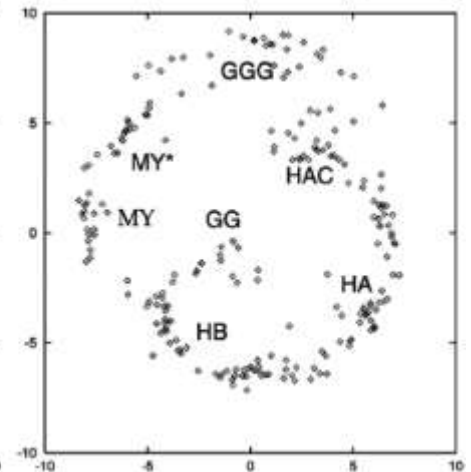


J_{ee}



J_{ff}

J_{ef}



MDS: Solution by Iterative Methods

- Common solution: Iterative optimization (e.g. gradient descent)
- Requires derivatives of objective function with respect to embedded point positions:

$$\begin{aligned}\frac{\partial J_{ee}}{\partial y_k} &= \frac{2}{\sum_{i<j} \delta_{ij}^2} \sum_{j \neq k} (d_{kj} - \delta_{kj}) \frac{y_k - y_j}{d_{kj}} \\ \frac{\partial J_{ff}}{\partial y_k} &= 2 \sum_{j \neq k} \frac{d_{kj} - \delta_{kj}}{\delta_{kj}^2} \cdot \frac{y_k - y_j}{d_{kj}} \\ \frac{\partial J_{ef}}{\partial y_k} &= \frac{2}{\sum_{i<j} \delta_{ij}} \sum_{j \neq k} \frac{d_{kj} - \delta_{kj}}{\delta_{kj}} \cdot \frac{y_k - y_j}{d_{kj}}\end{aligned}$$

Metric MDS: Solution via Kernel PCA

- Assume dissimilarity matrix $\Delta = (\delta_{ij})$ results from points in an unknown feature space:

$$\{\Phi_i \in \mathbb{R}^q, i = 1, \dots, n\}$$

- In this case:

$$\begin{aligned}\delta_{ij}^2 &= \|\Phi_i - \Phi_j\|^2 \\ &= \|\Phi_i\|^2 + \|\Phi_j\|^2 - 2\langle \Phi_i, \Phi_j \rangle\end{aligned}$$

- Idea:

If we can transform Δ into an inner product matrix, we can continue as in kernel PCA

Metric MDS: Reduction to Kernel PCA

- Inner product matrix of centered features Φ_i

$$K^c = [\langle \Phi_i, \Phi_j \rangle]$$

- Then, $\delta_{ij}^2 = K_{ii} + K_{jj} - 2K_{ij}$

Since features are centered $\sum_i \Phi_i = 0$,
we have that $\sum_i K_{ij} = 0$ and thus:

$$\frac{1}{N} \sum_i \delta_{ij}^2 = \frac{1}{N} \sum_i K_{ii} + K_{jj} \quad := \delta_{\oplus j}^2$$

$$\frac{1}{N} \sum_j \delta_{ij}^2 = K_{ii} + \frac{1}{N} \sum_j K_{jj} \quad := \delta_{i \oplus}^2$$


$$\frac{1}{N^2} \sum_i \sum_j \delta_{ij}^2 = \frac{2}{N} \sum_i K_{ii} \quad := \delta_{\oplus \oplus}^2$$

- Transform dissimilarities to

$$K_{ij} = -\frac{1}{2} (\delta_{ij}^2 - \delta_{i \oplus}^2 - \delta_{\oplus j}^2 + \delta_{\oplus \oplus}^2)$$

Metric MDS: Solution via Eigenvectors

1. Compute $K^c = -\frac{1}{2}HDH$ with $H = I - \frac{1}{N}\mathbf{1}\mathbf{1}^T$
 - If Δ is metric, this defines K^c such that $K^c = [\langle \Phi_i, \Phi_j \rangle]$
2. Compute spectral decomposition of K^c
 - $K^c = U\Lambda U^T$




3. Lower dimensional embedding:
 - Want k -dimensional coordinates
 - U_k = first k columns of U
 - Λ_k = $k \times k$ submatrix of Λ
 - Desired coordinates = $U_k \Lambda_k^{1/2}$

Properties of MDS

- Effort independent of dimensionality
 - Distance matrix is all that matters
- Solution unique only up to rigid transformation and reflection
- Approaches:
 - Optimization-based: Works with all matrices and different error measures
 - Analogous to kernel PCA: Assumes underlying metric space, permits projection of new points (if you know distances to all existing ones)

Example: Full-Brain Tractography

Problem: Nerve fiber pathways present themselves as an impenetrable knot of curves. How to interact with them (e.g., make a selection) in an easy way?



Idea: It is much simpler to interact with 2D views! Can't we represent each streamline as a point in 2D image space, such that similar streamlines are placed nearby?

Linked 2D Views




Image taken from [Jianu et al. 2009]

Section 4.3:


ISOMAP

Distances on Nonlinear Manifolds

- MDS allows us to use arbitrary distances
- On **nonlinear manifolds**, we should use **geodesic distances** (length of shortest connection on manifold), not Euclidean:




$$d(A,C) < d(A,B)$$



$$d(A,C) > d(A,B)$$


Idea: Estimating Geodesic Distance

- We are usually not given the manifold, only a set of discrete input points
- **Estimate geodesic distances:**
 - Use Euclidean distance for nearby points
 - Connect distant points by short hops between nearby points, measure length of shortest path



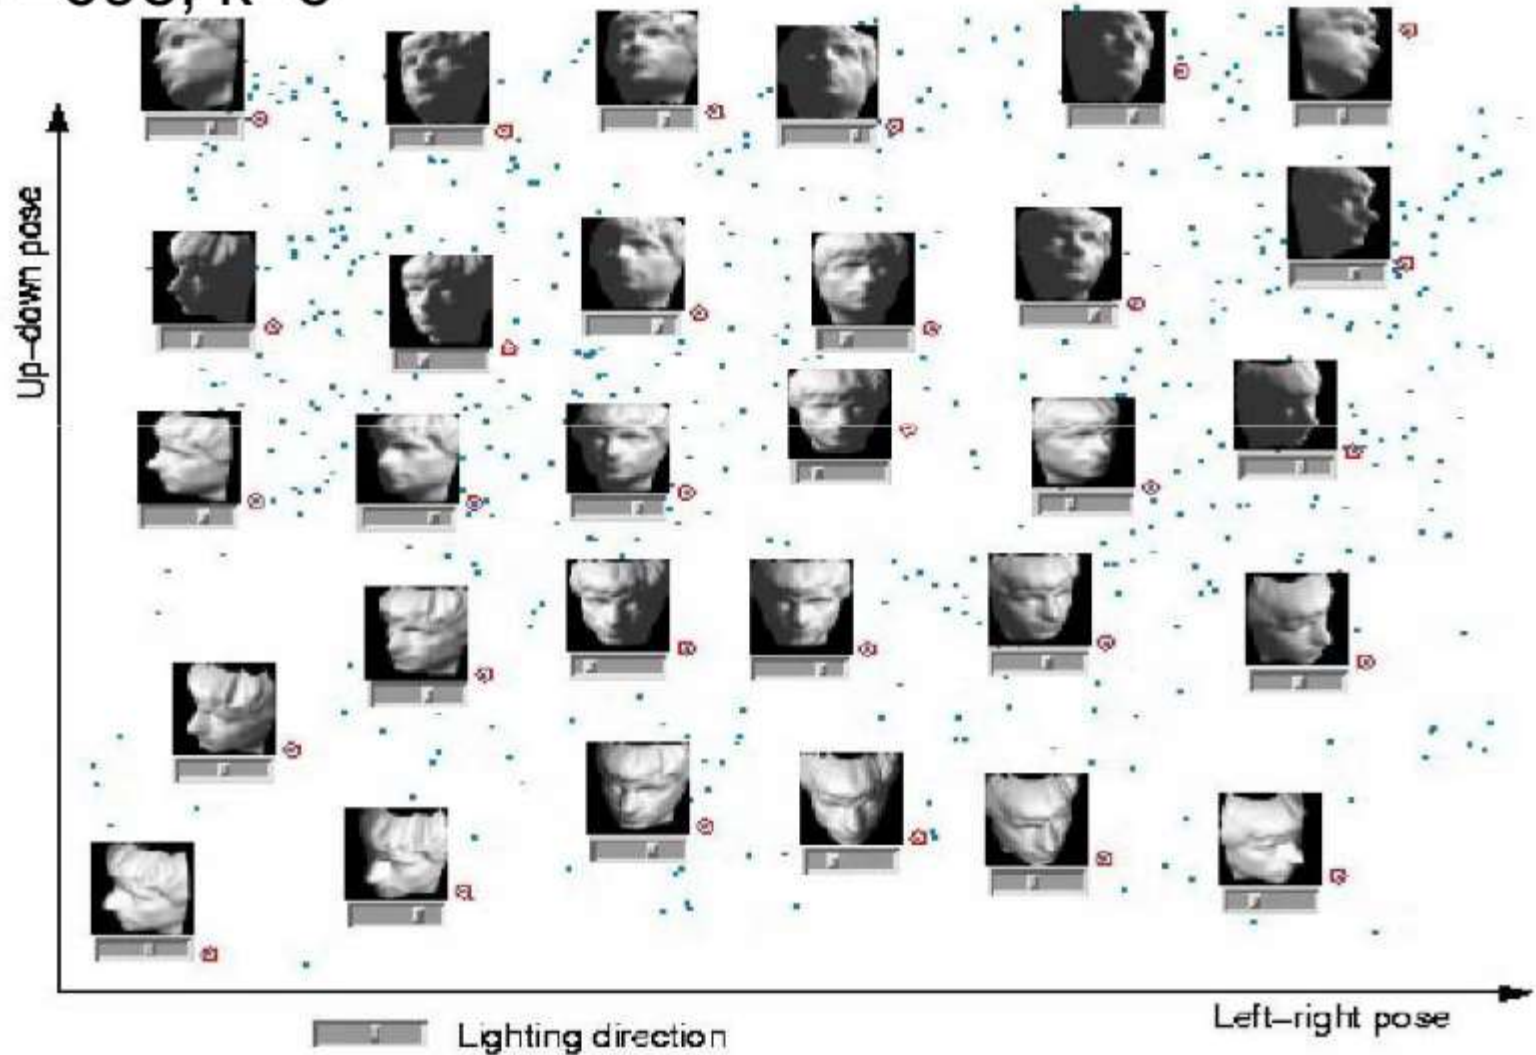
ISOMAP Algorithm

1. Build a neighborhood graph (e.g., k-NN) that approximates the manifold structure
 - Make sure it's connected
 - Very similar to spectral clustering
2. For the resulting graph, compute all-pairs shortest paths
3. Perform metric MDS based on resulting distance matrix



ISOMAP: Example result

$n = 698, k = 6$



Summary: Dimensionality Reduction

- Dimensionality reduction projects high-dimensional objects to low-dimensional space
 - **PCA:** Takes Euclidean input and finds linear projection that preserves maximum variance
 - **Kernel PCA:** Generalizes PCA to better preserve nonlinear structures
 - **MDS:** Can be applied even if original data does not have coordinates or if distances are non-Euclidean
 - **ISOMAP:** Combines MDS with distance measures on neighborhood graphs to approximate nonlinear manifolds